GATEWay: An Intelligent GATE DA Preparation Platform

With Advanced Recommendation Engines, Analytics, and Predictive Scoring

Khethavath Sunil Naik
Dosapati Ashirvadhan*
B.Tech Project Report

Abstract

GATEWay is a sophisticated, classroom-aware educational intelligence platform designed specifically for GATE (Graduate Aptitude Test in Engineering) Data Analytics (DA) preparation. The system provides dual-mode experiences: a comprehensive student workspace for structured practice, adaptive testing, personalized coaching, and performance tracking, coupled with an advanced teacher command center for classroom orchestration, assignment management, learner monitoring, and intervention planning.

The platform's core innovation lies in its unified architecture that seamlessly integrates multiple specialized recommendation engines, a telemetry-driven analytics layer powered by Supabase, graph-based adaptive sequencing, intelligent rule-based coaching, comprehensive performance analytics, and a novel predictive scoring model. This report provides a detailed technical exposition of the complete system architecture, mathematical formulas underlying key recommendations, all data models, recommendation engine implementations, teacher analytics pipelines, and the comprehensive GATE score prediction formula with detailed variable explanations. The document spans multiple dimensions: product vision, system design, algorithmic foundations, implementation details, and research assets, making it suitable for technical portfolios, advanced project vivas, research presentations, and academic evaluation.

Keywords

GATE DA, adaptive learning, recommendation systems, educational analytics, score prediction

1 Introduction and Problem Statement

1.1 Educational Challenges in GATE DA Preparation

The GATE (Graduate Aptitude Test in Engineering) Data Analytics examination is among India's most competitive postgraduate entrance examinations. Aspirants face several critical challenges:

(1) **Content Vastness**: The GATE DA syllabus spans mathematics, computer science fundamentals, probability, statistics, linear algebra, machine learning, and artificial intelligence—requiring diverse subject expertise.

(2) **Adaptive Learning Gap**: Traditional mock test portals provide static question sequences without personalizing difficulty or topic selection based on learner performance patterns.

(3) **Lack of Real-time Guidance**: Students complete tests in isolation without immediate, explanatory feedback or next-step recommendations tailored to their performance profile.

(4) **Teacher Engagement Inefficiency**: Educators lack systematic tools to identify struggling students, recommend targeted interventions, or track class-wide performance trends in real-time.

(5) **Performance Prediction Absence**: Students cannot reliably estimate their likely GATE score based on mock performance, hindering time allocation and preparation intensity decisions.

(6) **Incomplete Telemetry**: Existing platforms record only final scores, not the detailed event stream needed for explainability, remediation tracking, or longitudinal analytics.

1.2 Platform Vision

GATEWay addresses these challenges by providing:

- **Adaptive graph-based sequencing** that learns from each student's performance and dynamically adjusts question difficulty and topic selection.
- **Multi-engine recommendation portfolio** addressing next-best question selection, remediation paths, retry logic, suspicious rapid guessing detection, and personalized coaching.
- **Classroom orchestration tools** enabling teachers to author assignments, monitor submissions, identify at-risk students, and receive data-driven intervention suggestions.
- **Comprehensive telemetry** capturing both state (ELO, scores, progress) and events (actions, question answers, test completions) for rich observability.
- **Predictive scoring engine** that estimates likely GATE score bands based on learner proficiency factors, enabling informed preparation decisions.
- **Intelligent coaching system** using sophisticated rule-based logic to deliver personalized study recommendations and learning guidance.

The platform answers four fundamental educational questions:

(1) What should the student solve next?
(2) What should the student revise next?
(3) Which learners need teacher intervention?
(4) How can full learning paths be measured, reviewed, and explained?

(1)

2 System Architecture and Design Philosophy

2.1 Architecture Overview

The platform is designed as a layered, multi-component ecosystem shown in Figure 1.

Unpublished working draft. Not for distribution.

*Roll Number: 12341170

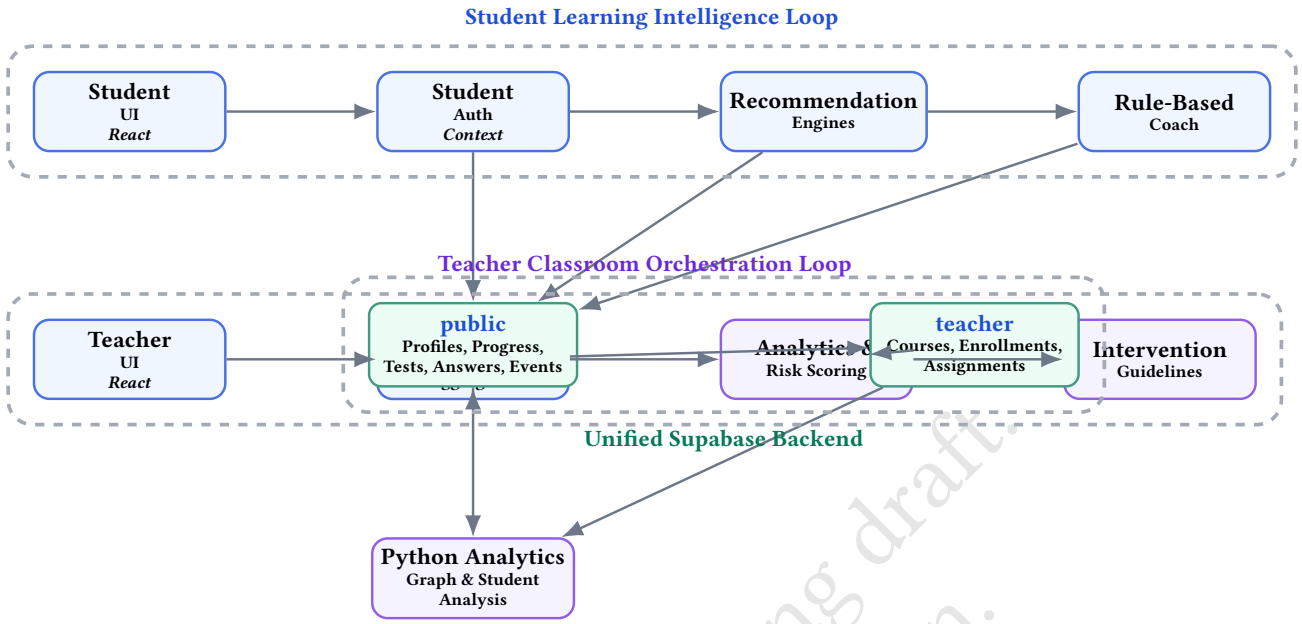

Figure 1: High-level system architecture. Student and teacher loops operate independently but converge on unified data schemas (public and teacher) stored in Supabase. Analytics notebooks enable exploratory research.

2.2 Technology Stack

2.3 Multi-Schema Data Modeling Philosophy

The architecture employs a single Supabase project with two exposed schemas, representing a critical design decision:

- **Public Schema**: Contains all student learning telemetry—profiles, progress, answered questions, test history, and activity events. Directly accessible to students, with teacher read access via row-level security.
- **Teacher Schema**: Contains classroom operations—teacher accounts, courses, student enrollments, assignments, and submissions. Isolated to teachers, with RLS preventing unauthorized access.

This separation:

- Maintains clear operational boundaries between student learning and teacher classroom management.
- Enables teacher analytics to traverse both schemas within row-level security constraints.
- Simplifies permission models by making the access pattern explicit.
- Supports efficient querying of class-level insights without mixing operational and learning data.

3 Data Model and Core Entities

3.1 Public Schema Tables

3.1.1 profiles. Stores authenticated user information, preferences, and account metadata.

3.1.2 user_progress. Tracks subject-level performance metrics for adaptive sequencing.

3.1.3 answered_questions. Records every question answer with full context for replay and analysis.

3.1.4 test_history. Aggregated test completion records with optional replay metadata.

3.1.5 activity_events. Event stream capturing all significant user actions for observability and analytics.

3.2 Teacher Schema Tables

3.2.1 teachers. Stores teacher account information and workspace metadata.

3.2.2 courses. Contains course information with unique join codes for student enrollment.

3.2.3 enrollments. Maps students to courses (many-to-many relationship).

3.2.4 assignments. Teacher-authored assignments with metadata and publication status.

3.2.5 submissions. Student assignment submissions with grades and feedback.

4 Student Learning Flow and Experience

4.1 Onboarding and Authentication

The student user flow begins with authentication via Supabase Auth:

(1) Student signs up with email and password.

Unpublished working draft. Not for distribution.

Layer	Technologies
Frontend Framework	React 18 with TypeScript for type-safe UI development
Styling	Tailwind CSS with custom design tokens for consistent theming
UI Components	Radix UI and shadcn/ui for accessible, headless components
Routing	React Router v6 for client-side routing
Icons	Lucide Icons for vector graphics and consistent iconography
State Management	React Context API with custom hooks for domain-specific state
Data Fetching	TanStack React Query for server state and caching
Backend	Supabase (PostgreSQL, Auth, REST API, Realtime, Edge Functions)
Database Schema	Single project with public and teacher schemas
Row-Level Security	Supabase RLS policies ensuring multi-tenancy isolation
Charting	Recharts for in-app visualizations, PGFPlots for research outputs
Coaching System	Deterministic rule-based coaching engine with custom heuristics
Testing	Vitest for unit tests, Testing Library for components, Playwright for E2E
Build Tools	Vite for fast bundling and development
Research	Python (NetworkX, Pandas, NumPy, Matplotlib) for analytics notebooks

Table 1: Comprehensive technology stack across all layers.

Field	Type	Purpose
user_id	UUID	Foreign key to Supabase Auth
full_name	Text	Display name
email	Text	Contact information
elo	Integer	Current ELO rating (starts at 1600)
tier	Text	Proficiency tier (e.g., Bronze, Silver, Gold, Platinum)
theme	Text	User theme preference (light, dark)
created_at	Timestamp	Account creation date
updated_at	Timestamp	Last profile update

Table 2: Profile table structure for storing user accounts and preferences.

Field	Type	Purpose
user_id	UUID	Student reference
subject	Text	Subject name (e.g., Linear Algebra, DBMS)
accuracy	Real	Fraction of correct answers in subject
questions_answered	Integer	Total questions attempted
weak_topic	Text	Identified weak area for remediation
updated_at	Timestamp	Last update

Table 3: User progress table tracking subject-level accuracy and weak areas.

(2) Supabase Auth creates a user account; a profiles record is automatically instantiated with default ELO (1600) and tier.

(3) Student logs in; the StudentAuthContext restores local and cloud state via React Query.

4.2 Enrollment and Course Joining

Students join teacher classrooms via join codes:

Field	Type	Purpose
id	UUID	Record identifier
user_id	UUID	Student reference
question_id	Text	Question identifier from bank
submitted_answer	Text	Student's selected or written answer
is_correct	Boolean	Correctness flag
time_spent	Integer	Seconds elapsed (for rapid-guess detection)
attempt_number	Integer	Attempt count for retry tracking
created_at	Timestamp	Answer submission time

Table 4: Answered questions table for complete answer history.

Field	Type	Purpose
id	UUID	Record identifier
user_id	UUID	Student reference
test_type	Text	Type (e.g., full_mock, topic_wise, adaptive)
test_name	Text	Human-readable test name
score	Integer	Raw score obtained
accuracy	Real	Fraction correct
completion_time	Integer	Total minutes spent
review_payload	JSON	Replayable data: question IDs, answers, metadata
created_at	Timestamp	Test completion time

Table 5: Test history table with optional review payload for explainability.

(1) Teacher creates a course and receives a unique 6-character join code.

(2) Student enters the code in the platform's "Join Course" modal.

Unpublished working draft. Not for distribution.

Field	Type	Purpose
id	UUID	Record identifier
user_id	UUID	Student reference
event_type	Text	Event category (e.g., question_answered, test_completed)
metadata	JSON	Event-specific context
created_at	Timestamp	Event occurrence time

Table 6: Activity events table for comprehensive telemetry and analytics.

Field	Type	Purpose
id	UUID	Unique course identifier
teacher_id	UUID	Course owner reference
name	Text	Course display name
join_code	Text	Unique enrollment code
description	Text	Course description
created_at	Timestamp	Course creation date

Table 7: Courses table for classroom creation and management.

(3) An enrollment record is created, linking the student to the course.

(4) From that point, student progress is visible to the teacher via the class roster.

4.3 Learning Modes

Students can engage with the platform through four primary learning pathways:

4.3.1 Full Mock Test Mode. Students attempt complete GATE DA mock papers covering all topics in a single, timed session (typically 3 hours).

Workflow:

- Student selects a mock paper from available options.
- The system randomizes question order within the paper.
- Timer enforces time limits and provides progress feedback.
- Upon completion, the system computes raw score, accuracy, and test metadata.
- Results are stored in test_history with a review_payload enabling answer replay.

4.3.2 Topic-Wise Practice Mode. Students practice questions restricted to a single topic (e.g., "Graph Algorithms," "Probability").

Workflow:

- Student selects a subject and topic.
- System retrieves all questions for that topic (ordered by ELO or randomized).
- Student solves questions one-by-one; each submission updates the answered_questions table.
- Performance within the topic updates the user_progress row for that subject.

4.3.3 Adaptive Practice Mode. The system uses the graph-based adaptive engine to recommend the next best question based on learner state.

Workflow:

- Student initializes an adaptive session.
- The **adaptive next-best question engine** queries the student's current ELO, momentum, weak topics, and answered question set.
- The engine navigates the question recommendation graph (see Section 5) and selects the next question with reasoning.
- Student submits their answer; ELO, subject accuracy, and answered questions are updated.
- The system computes the next recommendation and repeats until the student ends the session.

4.3.4 Assignment Mode. Teachers publish assignments; students access and submit responses.

Workflow:

- Teacher authors and publishes an assignment within a course.
- Students enrolled in that course see the assignment in their assignment list.
- Student attempts the assignment (typically a curated set of questions with time limits).
- Submission is recorded in the submissions table; teacher can grade and provide feedback.

4.4 Performance Tracking and Metrics

As students engage, the system tracks:

- **ELO Rating**: Updated after each question based on expected vs. actual performance (Elo chess rating system adapted for accuracy).
- **Subject Accuracy**: Fraction of correct answers within each subject (e.g., linear algebra accuracy = 74%).
- **Weak Topics**: Topics with below-threshold accuracy, flagged for remediation.
- **Test History**: Aggregated results from full mocks, topic-wise sessions, and adaptive practice.
- **Momentum**: A label (hot, steady, cold) computed from recent correctness streak.
- **Completion Rate**: Fraction of all available questions answered (used in score prediction unlock criteria).

5 Adaptive Graph-Based Recommendation Engine

5.1 Design Objectives and Philosophy

The adaptive engine is the intellectual core of GATEWay. Rather than employing simple difficulty-based sorting, it:

- Constructs a semantic recommendation graph over the entire question bank.

Unpublished working draft. Not for distribution.

- Models learner state as a multi-dimensional vector: current ELO, momentum, weak topics, recently attempted questions, and graph position.
- Navigates the graph using a hybrid scoring function combining heuristic rewards and graph-structure bonuses.
- Provides explainable recommendations with natural-language reasoning.
- Supports both standard adaptive sequences and targeted remediation after mistakes.

5.2 Question Graph Construction

The recommendation graph $G = (V, E)$ consists of:

- **Nodes (V)**: Question bank entries, each with attributes (subject, topic, ELO, marks, expected duration).
- **Edges (E)**: Directed edges representing pedagogical relationships.

5.2.1 *Edge Types and Semantics.*

(1) **Same-Topic Edges** (same-topic): Connect questions within the same subject and topic.
   - Useful for focused practice on a single skill.
   - Weight reflects difficulty progression: easier → harder or similar difficulty hops.

(2) **Subject-Flow Edges** (subject-flow): Connect topics within the same subject in logical progression.
   - Example: Linear Algebra 101 → Mathematical Properties → Eigenvalues and Eigenvectors.
   - Support topic-wise sequential learning.

(3) **Subject-Bridge Edges** (subject-bridge): Link related topics across different subjects.
   - Example: Probability (Statistics) ↔ Hypothesis Testing (Machine Learning).
   - Enable cross-domain skill transfer and integrated learning paths.

Edge weights encode:

$$w(u, v) = \alpha_1 \cdot d_{\text{ELO}} + \alpha_2 \cdot d_{\text{marks}} + \alpha_3 \cdot d_{\text{topic-proximity}} + \alpha_4 \cdot d_{\text{type-similarity}} \tag{2}$$

where d_* denotes normalized distances (lower is better), and α_i are empirically tuned weights.

5.3 Learner State Representation

At any point in an adaptive session, the learner state is represented as a tuple:

$$\mathbf{s} = (\text{ELO}, \text{momentum}, \text{weak_topics}, Q_{\text{answered}}, Q_{\text{served}}, \text{remediation_progress}) \tag{3}$$

where:

- ELO $\in [0, 3000]$: Current rating (higher = stronger).
- momentum $\in$ {hot, steady, cold}: Recent performance trend.
- weak_topics $\subseteq T$: Set of topics with accuracy $< \tau_{\text{weak}}$ (typically 60%).
- $Q_{\text{answered}} \subseteq V$: Set of questions already answered (never served again in a session).
- $Q_{\text{served}} \subseteq V$: Set of questions previously served (but not yet answered).
- remediation_progress: Metadata tracking remediation attempts for unresolved misses.

5.4 Target Policy Computation

The engine computes a target ELO and target difficulty:

5.4.1 *Target ELO.*

$$\text{ELO}_{\text{target}} = \begin{cases} \text{ELO} + 100 & \text{if momentum = hot and accuracy} \geq 0.8 \\ \text{ELO} + 50 & \text{if momentum = steady and accuracy} \geq 0.7 \\ \text{ELO} & \text{otherwise (steady or declining)} \\ \text{ELO} - 100 & \text{if momentum = cold and accuracy} < 0.5 \end{cases} \tag{4}$$

Intuitively: students on hot streaks are challenged with harder questions; struggling students receive easier review material.

5.4.2 *Target Difficulty.*

$$\text{difficulty}_{\text{target}} = \text{map_elo_to_difficulty}(\text{ELO}_{\text{target}}) \tag{5}$$

where the mapping is typically:

- ELO < 1600: Easy
- $1600 \leq$ ELO < 2000: Medium
- $2000 \leq$ ELO < 2500: Hard
- ELO ≥ 2500: Expert

5.5 Candidate Selection and Neighborhood Navigation

Given the learner state and target policy, the engine branches on whether an unresolved miss exists:

5.5.1 *Remediation Neighborhood.*

When a student answers a question incorrectly (unresolved miss), the engine defines a remediation neighborhood:

$$N_{\text{remed}}(q, h) = \{q' \in V : d_G(q, q') \leq h, \text{ not answered}, \text{ difficulty}(q') < \text{difficulty}(\tag{6}$$

where $d_G(q, q')$ is the graph distance (hops), and h (typically 2–3) is the maximum hop limit. This ensures:

- Questions are related to the originally missed question (nearby in graph).
- But easier (to build confidence before retry).
- And haven't been answered before.

5.5.2 *Current Question Neighborhood.*

In standard (non-remediation) adaptive flow:

$$N_{\text{current}}(q, k) = \{q' \in V : d_G(q, q') \leq k, \text{ not answered}, \text{ not served}\} \tag{7}$$

Typical $k = 2$ (2-hop neighbors).

Unpublished working draft. Not for distribution.

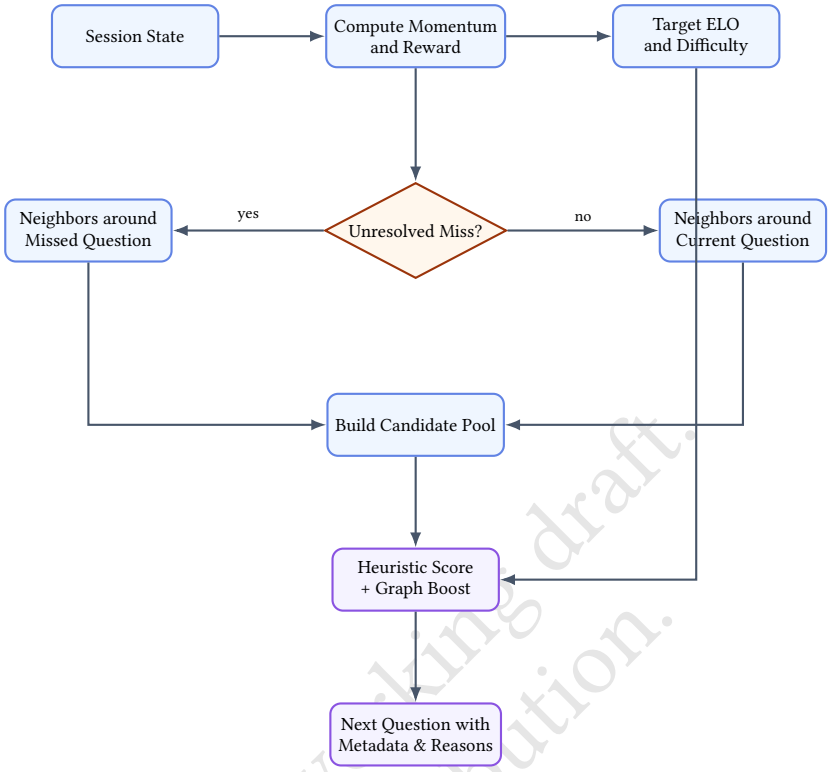

(a) Adaptive recommendation flow.

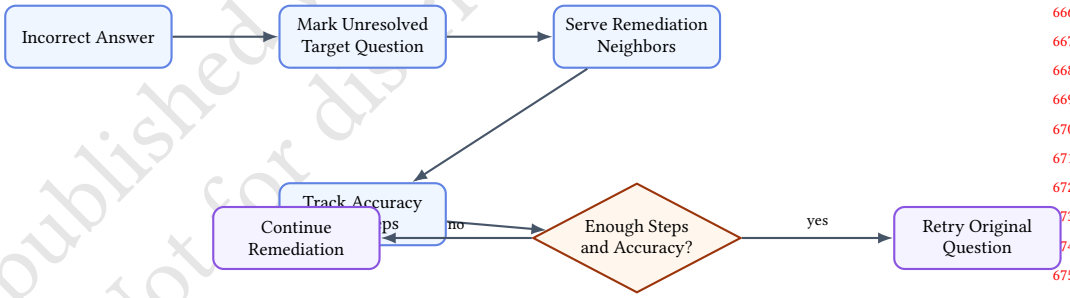

(b) Remediation and retry mechanism.

Unpublished working draft. Not for distribution.

Figure 2: Side-by-side visualizations of the adaptive recommendation pipeline and its remediation logic.

5.6 Candidate Scoring Function

After assembling the candidate pool C, the engine scores each candidate $q' \in C$:

$$\text{score}(q') = w_h \cdot h(q') + w_g \cdot g(q') \tag{8}$$

where:

- $h(q')$ is a **heuristic score** combining:
  (1) Distance to target difficulty: $d_{\text{diff}} = |\text{difficulty}(q') - \text{difficulty}_{\text{target}}|$
  (2) Distance to target ELO: $d_{\text{elo}} = |\text{ELO}(q') - \text{ELO}_{\text{target}}|$
  (3) Weak-topic bonus: if q' is in a weak topic, add bonus β_{weak}.
  (4) Type diversity: if q' type differs from the current question, add bonus β_{type}.

  Typically:

$$h(q') = 1.0 - \frac{w_1 \cdot d_{\text{diff}} + w_2 \cdot d_{\text{elo}}}{C_{\text{norm}}} + \beta_{\text{weak}} \cdot \mathbb{1}[q' \in \text{weak_topics}]$$

$$+ \beta_{\text{type}} \cdot \mathbb{1}[\text{type}(q') \neq \text{type}(q)] \tag{9}$$

- $g(q')$ is a **graph boost** rewarding structural properties:

$$g(q') = \gamma_1 \cdot \rho_{\text{in}}(q') + \gamma_2 \cdot \rho_{\text{out}}(q') + \gamma_3 \cdot c(q') \tag{10}$$

  where:

  – $\rho_{\text{in}}(q')$ = normalized in-degree (how many questions lead to q').

– $\rho_{\text{out}}(q')$ = normalized out-degree (how many questions q' leads to).
– $c(q')$ = normalized centrality measure (e.g., betweenness centrality).

Nodes with high centrality are pedagogically important hub questions, reflecting well-established learning pathways.

Weights w_h, w_g, and coefficients are tuned empirically via A/B testing or offline replay.

5.7 Output and Explainability

The engine returns the top-k (typically $k = 1$) candidate with:

- Question ID and full metadata (text, options, marks, expected difficulty).
- Natural-language explanation: e.g., "Continuing on Linear Algebra—the next question is harder to strengthen your eigenvalue skills."
- Graph metadata: hop distance, centrality signal, neighborhood type (remediation vs. standard).
- Target ELO and expected difficulty for transparency.

6 Remediation and Retry Engine

When a student submits an incorrect answer, the system marks the question as an unresolved miss. The remediation engine manages the recovery process:

6.1 Remediation Workflow

(1) Flag question as unresolved miss
(2) Serve remediation neighbors (easier, related)
(3) Track accuracy and step count within remediation sequence
(4) After N_{remed} successful steps or accuracy $\geq A_{\text{remed}}$: unlock retry of original
(5) If retry is correct: clear miss, mark as resolved
(6) If retry is incorrect: re-queue for remediation

(11)

6.2 Remediation Exit Criteria

Remediation concludes when:

$$\text{steps_taken} \geq N_{\text{remed}} \quad \wedge \quad \text{remediation_accuracy} \geq A_{\text{remed}} \tag{12}$$

Typical values: $N_{\text{remed}} = 3$ remediation steps, $A_{\text{remed}} = 0.75$ (75% accuracy).

If these criteria are not met within a session, the miss carries forward; the retry is queued for the next session.

7 Practice Review and Rapid-Guess Detection Engine

7.1 Motivation: Quality of Evidence

Not all correct answers represent equal mastery. A student might answer a complex multi-step problem correctly but in 10 seconds—suspicious and likely guessing. The rapid-guess detection engine in src/lib/practiceReview.ts flags such scenarios.

7.2 Threshold Computation

For each question, the engine estimates a base minimum solve time:

$$t_{\min}(q) = \alpha_0 + \alpha_1 \cdot \text{stem_length}(q) + \alpha_2 \cdot \text{options_length}(q) + \alpha_3 \cdot \text{complexity_index}(q) \tag{13}$$

where complexity index is higher for questions labeled as multi-step or requiring computation.

Typical coefficient values (empirically tuned):

- $\alpha_0 = 10$ seconds (baseline reading time).
- $\alpha_1 = 0.02$ seconds per character in stem.
- $\alpha_2 = 0.01$ seconds per character in option set.
- $\alpha_3 = 30$ seconds for high-complexity questions.

7.3 Penalty Application

If $t_{\text{spent}}(q) < t_{\min}(q)$:

$$\text{ELO_penalty} = (1 - \frac{t_{\text{spent}}}{t_{\min}}) \cdot 20 \text{ points} \tag{14}$$

$$\text{ELO_gain}(q) = \max(0, \text{standard_gain} - \text{ELO_penalty}) \tag{15}$$

The student still receives marks, but their ELO gain is reduced, preventing rapid guessing from inflating proficiency ratings.

8 Coach Engine

8.1 Purpose and Scope

The coach provides personalized study recommendations, learning plans, and guidance to students. Rather than being a chatbot, it focuses on structured, actionable coaching:

- Identifies weak and strong subject areas.
- Recommends topic prioritization and revision strategies.
- Generates same-day, 3-day, and 7-day study plans.
- Provides mock preparation strategies and effort allocation advice.
- Responds to student questions with carefully crafted rule-based recommendations.

8.2 Rule-Based Implementation Architecture

The coach is built on a pure rule-based heuristic engine with no external dependencies. All coaching logic is hand-crafted and parameterized based on empirical observations from GATE DA preparation patterns.

Key Properties:

(1) **Fast Execution**: Response generation completes in milliseconds, ensuring responsive user experience.
(2) **Deterministic and Explainable**: Every recommendation follows transparent, understandable rules that can be audited and explained to students.
(3) **No External Dependencies**: The coach operates offline and does not require API calls, network connectivity, or external services.
(4) **Custom Tuning**: Parameters are adjusted based on observed student behavior patterns within the platform.

Unpublished working draft. Not for distribution.

Condition	Generated suggestion
Weak subjects available	Focus on w_1 today; you still have $(100 - \text{acc}_{w_1})$% room for improvement.
Streak > 5	Great momentum. Tackle harder problems in s_1 today.
Otherwise	Review fundamentals in w_1 to build confidence.

Table 8: Daily study suggestion rules used by the deterministic coach.

8.3 Coach Input State

The coach receives:

$$\mathbf{c} = (\text{ELO}, \text{tier}, \text{overall_acc}, \text{weak_subjects}, \text{strong_subjects}, \text{streak}, \text{recent_tests}, \text{focus_area}) \tag{16}$$

where:

- ELO $\in [0, 3000]$: Current proficiency.
- tier: Categorical proficiency level (Bronze, Silver, Gold, Platinum).
- overall_acc $\in [0, 1]$: Fraction of all questions answered correctly.
- weak_subjects: Set of subjects with accuracy $< \tau_{\text{weak}}$.
- strong_subjects: Set of subjects with accuracy $\geq \tau_{\text{strong}}$.
- streak: Current correct-answer streak length.
- recent_tests: List of recent test results (score, accuracy, time).
- focus_area: Student-specified topic of interest or concern.

8.4 Deterministic Coach Logic

The fallback coach implements heuristic rules:

8.4.1 Daily Study Suggestion.

8.4.2 7-Day Study Plan. The coach allocates effort across weak, strong, and new topics:

$$\boxed{\begin{array}{ll} \text{Day 1-2 : Weak topic 1} & (40\% \text{ effort}) \\ \text{Day 3-4 : Weak topic 2} & (30\% \text{ effort}) \\ \text{Day 5-6 : Strong topic (mock)} & (20\% \text{ effort}) \\ \text{Day 7 : Revision or new topic} & (10\% \text{ effort}) \end{array}} \tag{17}$$

Effort allocation is parameterized by tier and recent performance.

8.5 Coach Reasoning Logic

The rule-based engine employs a multi-step pipeline to generate recommendations:

8.5.1 Step 1: Student Profile Analysis. The system analyzes the student input state to compute:

- Proficiency level based on current ELO (low: <1800, medium: 1800-2200, high: >2200).
- Performance stability (comparing recent vs. historical accuracy).
- Effort allocation efficiency (questions per topic studied).
- Motivation signal (streak length and session frequency).

8.5.2 Step 2: Weakness Severity Ranking. Weak subjects are ranked by severity:

$$\text{severity}(s) = (1 - \text{accuracy}(s)) \times \text{weight}(s) \tag{18}$$

where weight reflects importance in GATE DA (e.g., probability gets higher weight than niche topics).

8.5.3 Step 3: Personalized Recommendations. Based on the student profile, the coach selects from a curated set of advice templates:

- If ELO < 1800 and accuracy < 50%: Focus on fundamentals and concept clarity.
- If ELO $\in [1800, 2200]$ and improving: Maintain momentum by tackling harder questions.
- If ELO > 2200 and accuracy > 75%: Shift focus to speed optimization and mock practice.
- If declining trend detected: Recommend topic-specific review and reduced difficulty.

8.5.4 Step 4: Study Plan Generation. The 7-day plan is constructed by:

- Day 1-2: Highest-severity weak topic (primary focus).
- Day 3: Second-highest-severity weak topic and foundational concepts.
- Day 4-5: Integrated practice covering multiple topics.
- Day 6: Mock test or full-paper attempt to assess integrated readiness.
- Day 7: Targeted revision of mistakes from the mock test.

Each day's plan includes specific exercise counts, difficulty targets, and expected study duration.

Figure 3: Adaptive Learning Graph showing question relationships and student progression

9 Teacher Analytics and Intervention Intelligence

9.1 Teacher Workspace Aggregation

The teacher workspace hook (in hooks/useTeacherWorkspace.ts) aggregates classroom data:

(1) Fetch all teaching credentials for the logged-in teacher.
(2) Query all courses created by the teacher.
(3) For each course, fetch enrollments, assignments, and submissions.

Unpublished working draft. Not for distribution.

Metric	Definition
Overall Accuracy	acc = $\frac{\#\text{ correct answers}}{\#\text{ total answered}}$
ELO	Current rating from profiles
Weak Topics (top 3)	Subjects with accuracy < 0.6, sorted by severity
Assignment Completion	Fraction of assigned questions submitted
Risk Level	Composite score (see Section 9.3)
Recent Activity	Last action timestamp (question answered, test completed, etc.)
Performance Trend	7-day accuracy comparison

Table 9: Student summary metrics computed for teacher viewing.

(4) Execute a joined query over public and teacher schemas to retrieve student profiles, progress, test history, and activity events for enrolled learners.

This creates a unified view of classroom state available to the teacher dashboard.

9.2 Teacher Summary Computations

9.2.1 Student Summaries. For each enrolled student, compute:

9.2.2 Course Summaries. For each course, compute:

- Enrollment count: number of enrolled students.
- Average accuracy: mean of all student accuracies.
- Completion rate: fraction of assigned questions completed.
- Weak topics (class-wide): frequency of below-threshold accuracy by subject.
- At-risk student count: students with risk level $> T_{\text{risk}}$ (typically 0.7).
- Top performers: students with accuracy $> T_{\text{top}}$ (typically 0.8).

9.3 Risk Scoring

Risk distills learner challenges into a single actionable score. The composite risk score is:

$$\text{Risk}(s) = w_1 \cdot r_{\text{acc}}(s) + w_2 \cdot r_{\text{comp}}(s) + w_3 \cdot r_{\text{elo}}(s) \tag{19}$$

where:

- $r_{\text{acc}}(s) = 1 - \text{accuracy}(s)$: Accuracy-based risk (lower accuracy = higher risk).
- $r_{\text{comp}}(s) = 1 - \text{completion_rate}(s)$: Completion-based risk (not engaging enough).
- $r_{\text{elo}}(s) = \begin{cases} 1 & \text{if ELO}(s) < 1800 \\ 0 & \text{otherwise} \end{cases}$: ELO readiness (students under threshold need support).

Typical weights: $w_1 = 0.5$, $w_2 = 0.3$, $w_3 = 0.2$ (accuracy is the primary risk signal).

Students with Risk ≥ 0.7 appear in the teacher's intervention queue.

9.4 Intervention Recommendations

For each at-risk student, the teacher analytics module recommends actions:

(1) **Identify the top weak topic** with lowest accuracy.
(2) **Suggest a targeted assignment**: author and publish an exercise within that weak topic.
(3) **Check remediation progress**: query answered_questions to see if the student is successfully remediating past misses.
(4) **Schedule a follow-up**: recommend a check-in date based on the assignment deadline.

Example recommendation output:

```
Recommendation: Priya needs help with Probability &
Statistics.
Current accuracy: 45%
Recent miss: Q2834 (binomial distribution)
Suggested action: Publish a 5-question probability
assignment.
Follow-up: Check progress in 2 days
```

10 GATE Score Prediction Engine

10.1 Vision and Purpose

The GATE score prediction module src/components/GateScorePrediction.tsx addresses a critical learner need: estimating likely final GATE scores based on current preparation progress. This enables students to:

- Set realistic targets (e.g., target 85+ percentile rank).
- Allocate remaining study time effectively.
- Identify whether their preparation pace is sufficient.
- Build confidence through informed end-state projections.

10.2 Unlock Criteria

The score predictor is deliberately kept exclusive to avoid premature or unreliable predictions. It unlocks only when:

$$\text{Predictor UNLOCKED} \iff \begin{cases} \text{ELO} \geq 2500 \quad \text{AND} \\ \text{full_mock_completion} \geq 0.70 \quad \text{AND} \\ \text{topic_wise_completion} \geq 0.70 \quad \text{AND} \\ \text{adaptive_completion} \geq 0.70 \end{cases} \tag{20}$$

This ensures the model operates on substantial evidence:

- ELO 2500+ indicates solid proficiency.
- 70% completion across all three learning modes ensures broad coverage and engagement depth.

10.3 Core Formula: Predicted Score Basis

Once unlocked, the predictor generates a **score basis**—a normalized composite score from four key dimensions:

$$\text{ScoreBasis} = w_1 \cdot \text{Acc}_{\text{norm}} + w_2 \cdot \text{ELO}_{\text{norm}} + w_3 \cdot \text{Cons}_{\text{norm}} + w_4 \cdot \text{Impr}_{\text{norm}} \tag{21}$$

where the weights are: $w_1 = 0.40$, $w_2 = 0.30$, $w_3 = 0.20$, $w_4 = 0.10$.

Below, we explain each component in detail.

Unpublished working draft. Not for distribution.

10.3.1 Component 1: Overall Accuracy (Acc). **Conceptual Definition**:

$$\text{Accuracy} = \frac{\#\text{ questions answered correctly}}{\#\text{ total questions answered}} \tag{22}$$

Variable Definitions:

- C: Count of correct answers (from answered_questions where is_correct=true).
- T: Total questions answered.
- $\text{Accuracy}_{raw} = \frac{C}{T}$.

Normalization:

$$\text{Acc}_{\text{norm}} = \min\left(1.0, \frac{\text{Accuracy}_{raw} - 0.5}{0.35}\right) \tag{23}$$

This maps raw accuracy to $[0, 1]$:

- Accuracy $\leq 50\%$ (struggling) $\rightarrow$ normalized score ≤ 0.
- Accuracy $= 85\%$ (strong) $\rightarrow$ normalized score ≈ 1.0.
- Accuracy $> 85\%$ $\rightarrow$ capped at 1.0.

The lower bound (50%) reflects that random guessing on MCQC with 4 options yields 25% accuracy; 50% is a meaningful threshold above guessing.

10.3.2 Component 2: ELO Rating (ELO_{norm}). **Conceptual Definition**: The ELO rating (from profiles) reflects learner proficiency through wins and losses relative to question difficulty. Higher ELO = stronger learner.

Variable Definitions:

- $\text{ELO}_{current}$: Current ELO from user profile (range: 0—3000).
- $\text{ELO}_{baseline} = 1600$: Starting ELO.
- $\text{ELO}_{max} = 2800$: Typical maximum reached by very strong learners.

Normalization:

$$\text{ELO}_{\text{norm}} = \frac{\text{ELO}_{current} - \text{ELO}_{baseline}}{\text{ELO}_{max} - \text{ELO}_{baseline}} = \frac{\text{ELO}_{current} - 1600}{1200} \tag{24}$$

This maps:

- ELO = 1600 (baseline) $\rightarrow$ 0.0.
- ELO = 2400 (strong) $\rightarrow$ 0.667.
- ELO = 2800 (very strong) $\rightarrow$ 1.0.

10.3.3 Component 3: Subject Consistency (Cons_{norm}). **Conceptual Definition**: A learner with wildly varying subject accuracies (e.g., 90% in one subject, 40% in another) is less prepared than one with balanced proficiency. Consistency measures how evenly performance is distributed across subjects.

Variable Definitions:

- $S = \{s_1, s_2, \ldots, s_n\}$: Set of subjects (e.g., Linear Algebra, DBMS, Probability, ML, AI).
- acc_i = accuracy in subject s_i (from user_progress).
- $\bar{\text{acc}} = \frac{1}{n}\sum_{i=1}^{n} \text{acc}_i$: Mean accuracy across subjects.

Consistency Metric (Inverse Coefficient of Variation):

$$\text{Consistency}_{raw} = 1 - \frac{\sigma(\text{acc}_1, \ldots, \text{acc}_n)}{\bar{\text{acc}} + \epsilon} \tag{25}$$

where σ is standard deviation and ϵ is a small value (e.g., 0.01) to avoid division by zero.

Intuitively:

- If all subjects have accuracy 75%: $\sigma = 0 \Rightarrow$ Consistency = 1.0 (perfect).
- If accuracies are [90%, 40%, 70%, 65%, 60%], σ is large $\Rightarrow$ Consistency is lower.

Normalization:

$$\text{Cons}_{\text{norm}} = \max(0, \text{Consistency}_{raw}) \tag{26}$$

Bounded to $[0, 1]$.

Interpretation: A student with 80% overall accuracy will score higher if that 80% is evenly distributed (75%, 85%, 78%, 82%) vs. skewed (95%, 45%, 80%, 85%). The former shows better mastery of breadth; the latter shows gaps that may cause score loss in the actual exam.

10.3.4 Component 4: Improvement Factor (Impr_{norm}). **Conceptual Definition**: Students showing upward performance trends (improving over time) are likely to maintain or extend that improvement. The improvement factor rewards recent accelerating progress.

Variable Definitions:

- $T_{\text{recent}} \in (0, T_{\text{all}}]$: Recent time window (e.g., last 14 days).
- $\text{acc}_{\text{recent}}$: Accuracy in recent window.
- $\text{acc}_{\text{prior}}$: Accuracy in all prior data.
- $\Delta\text{acc} = \text{acc}_{\text{recent}} - \text{acc}_{\text{prior}}$: Improvement magnitude.

Improvement Metric:

$$\text{Improvement}_{raw} = \begin{cases} \min(0.15, \Delta\text{acc}) & \text{if } \Delta\text{acc} > 0 \quad \text{(improving)} \\ \max(-0.15, \Delta\text{acc}) & \text{if } \Delta\text{acc} \leq 0 \quad \text{(declining)} \end{cases} \tag{27}$$

This captures improvements and declines up to ±15 percentage points.

Normalization:

$$\text{Impr}_{\text{norm}} = \frac{\text{Improvement}_{raw} + 0.15}{0.30} \tag{28}$$

Maps to $[0, 1]$:

- $\Delta\text{acc} = -15\%$ (declining) $\rightarrow$ 0.0.
- $\Delta\text{acc} = 0\%$ (stable) $\rightarrow$ 0.5.
- $\Delta\text{acc} = +15\%$ (improving) $\rightarrow$ 1.0.

Interpretation: A student improving by 5 percentage points recently is rewarded; one declining is penalized. This captures momentum.

10.4 Score Band Generation from Basis

Once ScoreBasis $\in [0, 1]$ is computed, the predictor generates GATE score estimates. The GATE DA exam scoring is typically:

- Total marks: 100 (occasionally 85 or 90 depending on exam).
- Question distribution: usually 55 MCQ (1 mark each) + 10 NAT (numerical) questions (various marks).

For simplicity, assume 100-mark scale.

Unpublished working draft. Not for distribution.

Expected score band	Generated guidance
Expected ≥ 75	Excellent. Expect a GATE score around 80–90 and focus on maintaining consistency across all topics.
$60 \leq$ Expected < 75	Good progress. Expect a score around 60–70, strengthen weak topics, and increase mock-test frequency.
Expected < 60	Preparation is on track for 40–60. Increase study intensity, prioritise weak areas, and attempt one mock test per week.

Table 10: Guidance text produced by the score predictor.

10.4.1 Expected Score.

$$\text{Expected_Score} = 40 + 50 \cdot \text{ScoreBasis} \tag{29}$$

This maps ScoreBasis to an expected GATE score:

- ScoreBasis = 0.0 (barely qualifying) $\longrightarrow$ Expected = 40.
- ScoreBasis = 0.5 (moderate) $\longrightarrow$ Expected = 65.
- ScoreBasis = 1.0 (excellent) $\longrightarrow$ Expected = 90.

10.4.2 Minimum and Maximum Score Bands. To provide a confidence range, the model estimates bounds:

$$\text{Minimum_Score} = \text{Expected_Score} - 10 \cdot (1 - \text{ScoreBasis}) \tag{30}$$

$$\text{Maximum_Score} = \text{Expected_Score} + 10 \cdot \text{ScoreBasis} \tag{31}$$

Intuitively:

- High ScoreBasis (near 1.0) $\longrightarrow$ tight band (high confidence).
- Low ScoreBasis (near 0.0) $\longrightarrow$ wider band (more uncertainty).

Example:

- ScoreBasis = 0.8 $\Longrightarrow$ Expected = 80, Min = 72, Max = 88.
- ScoreBasis = 0.4 $\Longrightarrow$ Expected = 60, Min = 46, Max = 64.

10.4.3 Confidence Score. Confidence reflects certainty in the prediction:

$$\text{Confidence} = 0.5 + 0.4 \cdot \text{ScoreBasis} + 0.1 \cdot \text{Completion_Rate} \tag{32}$$

where Completion_Rate is the fraction of all GATE DA questions answered (0 to 1).

This maps to:

- Strong basis + high completion $\longrightarrow$ Confidence near 1.0 (very confident).
- Weak basis + low completion $\longrightarrow$ Confidence near 0.5 (moderate confidence).

10.4.4 Recommendation Text. The predictor generates contextual guidance:

10.5 Worked Example

Assume a student with:

- Accuracy: 78%.
- ELO: 2350.
- Subject accuracies: [85%, 72%, 70%, 75%, 80%] (standard deviation: 6%).
- Mean: 76.4%.
- Consistency: $1.0 - \frac{0.06}{0.764+0.01} \approx 0.922$.
- Recent accuracy (last 14 days): 82%; all prior: 75%; Δacc = 7%.

Normalization:

- $\text{Acc}_{\text{norm}} = \min(1.0, \frac{0.78-0.5}{0.35}) = \frac{0.28}{0.35} \approx 0.80$.
- $\text{ELO}_{\text{norm}} = \frac{2350-1600}{1200} \approx 0.625$.
- $\text{Cons}_{\text{norm}} = 0.922 \approx 0.92$.
- $\text{Impr}_{\text{norm}} = \frac{0.07+0.15}{0.30} \approx 0.73$.

Score Basis:

$$\text{ScoreBasis} = 0.40 \cdot 0.80 + 0.30 \cdot 0.625 + 0.20 \cdot 0.92 + 0.10 \cdot 0.73 \tag{33}$$

$$= 0.32 + 0.1875 + 0.184 + 0.073 = 0.7645 \tag{34}$$

Score Predictions:

- Expected: $40 + 50 \cdot 0.7645 = 40 + 38.225 = 78.2 \approx 78$.
- Minimum: $78 - 10 \cdot (1 - 0.7645) = 78 - 2.355 = 75.6 \approx 76$.
- Maximum: $78 + 10 \cdot 0.7645 = 78 + 7.645 = 85.6 \approx 86$.
- Confidence: $0.5 + 0.4 \cdot 0.7645 + 0.1 \cdot 0.70 \approx 0.5 + 0.306 + 0.07 = 0.876$ (high confidence).

Recommendation:

> Good progress. Expected GATE score: 76–86 (best estimate: 78). Consistency is strong across topics. Recent acceleration (+7%) is encouraging, so maintain this pace with daily practice on weak areas.

11 Telemetry, Persistence, and Event-Driven Observability

11.1 State Persistence Strategy

The platform persists learner state in two forms:

11.1.1 Persistent State (Tables).

- ELO, tier, and profile metadata in profiles.
- Subject accuracy and weak-topic tracking in user_progress.
- Complete answer history in answered_questions (for replay and remediation).
- Test summaries with optional review_payload in test_history.

This state is queried by the recommendation engines, analytics modules, and student dashboards.

11.1.2 Event Stream (activity_events table). In addition to state, every significant action is logged as an event:

11.2 Why This Design Matters

The dual stream (state + events) is critical:

(1) **Explainability**: Given a student's current ELO, the event stream enables reconstruction of how they arrived at that rating (what answers led to ELO changes).
(2) **Observability**: Teachers and administrators can inspect recent activity, understand student engagement patterns, and identify behavioral anomalies.
(3) **Research and Analytics**: The event stream enables longitudinal studies, cohort analysis, and identification of effective learning sequences.

Unpublished working draft. Not for distribution.

Event Type	Metadata
user_signed_up	user_id, email, timestamp
user_signed_in	user_id, timestamp
question_answered	user_id, question_id, is_correct, time_spent
test_started	user_id, test_type, test_name
test_completed	user_id, test_type, score, accuracy
graph_path_started	user_id, session_id
graph_path_completed	user_id, session_id, questions_count, accuracy
course_joined	user_id, course_id
course_left	user_id, course_id
assignment_published	teacher_id, course_id, assignment_id
submission_received	student_id, assignment_id, score

Table 11: Sample activity events and associated metadata.

(4) **Audit Trail**: Every significant action is recorded, supporting academic integrity monitoring.

(5) **Dashboard and History Views**: Students can replay past sessions using the review_payload, understanding which questions they missed and why.

12 History Review and Explainable Replaying

A salient feature is the **review_payload** field in test_history. When a test completes, the system optionally records:

- List of question IDs presented (in order).
- For each question: submitted answer, correctness, time spent, any rapid-guess warning.
- ELO change due to the question (if available).
- Explanation metadata (graph information, remediation notes).

This enables a **replayable history view**:

(1) Student selects a past test from history.

(2) System reconstructs the session question-by-question.

(3) For each question, the app displays: question text, student's answer, correct answer, explanation.

(4) Rapid-guess warnings are highlighted: "This was answered in 5 seconds, but estimated minimum was 12 seconds. ELO boost was reduced."

This transforms history from opaque score summaries into interactive learning artifacts.

13 Research Assets and Analytical Notebooks

The repository includes Python-based analysis tools:

- generate_gate_da_question_graph_notebook.py: Creates a Jupyter notebook analyzing the question recommendation graph (centrality, connectivity, degree distribution).
- generate_gate_da_student_progress_notebook.py: Generates a notebook with student cohort analytics (ELO distribution, accuracy trends, subject performance heatmaps).
- Teacher dashboard mockup generator: Creates sample visualizations showing class-wide metrics.

These notebooks use the same data model and provide exploratory analysis complementary to the live UI.

14 Testing and Quality Assurance

14.1 Test Coverage for Adaptive Engine

The adaptive engine is the highest-risk component, so it receives focused testing:

(1) **Same-Topic Progression**: After answering a question correctly, the next recommendation should be in the same topic but harder (or similar difficulty).

(2) **Cross-Topic Remediation**: After answering incorrectly, remediation recommendations should be easier and preferably in the same topic.

(3) **Retry Logic**: After sufficient remediation with high accuracy, a retry of the missed question should be queued.

(4) **Exclusion of Answered Questions**: Previously answered questions must never be served again in a session.

(5) **ELO Bounds**: ELO changes must stay within expected ranges, and rapid-guess penalties must be correctly applied.

Tests are written in Vitest with fixtures for sample graphs and learner states.

14.2 Component Testing

React components (especially recommendation displays and adaptive flows) are tested using Testing Library to ensure:

- Recommendations are rendered with rationale text.
- Answer submission correctly updates state.
- History replay correctly reconstructs past sessions.
- Teacher dashboards correctly filter and display student data.

14.3 E2E Testing

Playwright tests cover end-to-end flows:

(1) Student signup, login, and profile creation.

(2) Joining a teacher course via code.

(3) Attempting an adaptive session and verifying ELO updates.

(4) Reviewing history and replaying answers.

(5) Teacher viewing classroom dashboards and intervention recommendations.

15 Architectural Strengths

(1) **Clean Role Separation**: Student and teacher experiences are conceptually distinct, reflected in separate UI paths and data schemas.

(2) **Explainable Recommendations**: Every recommendation includes reasoning (graph hop distance, difficulty target, weak-topic signal), supporting learner trust and educator understanding.

Unpublished working draft. Not for distribution.

(3) **Comprehensive Telemetry**: The activity_events stream enables observability into learner behavior, supporting debugging, analytics, and curiosity-driven research.

(4) **Intelligent Rule-Based Coaching**: The coaching engine is built on carefully designed heuristics and domain knowledge, ensuring transparent, understandable, and reliable recommendations without external dependencies.

(5) **Replayable History**: The review_payload upgrade converts static score records into interactive learning artifacts for deeper reflection.

(6) **Remediation-Aware Sequencing**: The adaptive engine doesn't just select harder questions; it proactively identifies unresolved misses and queues targeted recovery sequences.

(7) **Teacher-Centric Actional Insights**: Analytics output prioritizes actionable recommendations (suggest specific assignments, identify at-risk students) over passive reporting.

(8) **Thoughtful Score Prediction**: The GATE score predictor balances mathematical rigor (detailed formula with four normalized components) with practical utility (confidence bands, recommendation text).

(9) **Research Layer**: Python notebooks and analysis tooling complement the live UI, supporting exploratory analytics and academic inquiry.

16 Limitations and Future Enhancements

16.1 Current Limitations

(1) **Formal Experiment Tracking**: The platform lacks a first-class A/B testing framework. Recommendation policy changes are not yet systematically compared via randomized experiments.

(2) **Longitudinal In-App Analytics**: While the activity_events stream is rich, fewer longitudinal visualizations are surfaced directly in the app. Most longitudinal analysis happens in notebooks.

(3) **Score Prediction History**: Current predictions are point-in-time. Persisting historical predictions would enable trend analysis ("How did my predicted score change over the past month?").

(4) **Spaced Repetition**: The remediation engine is reactive (triggered after misses). A proactive spaced-repetition system would queue forgotten topics for revisiting.

(5) **Prerequisite Graphs**: The current graph is similarity-based. An explicit prerequisite graph (e.g., Linear Algebra Fundamentals → Eigenvalues) would improve sequencing.

(6) **Cross-Subject Integration**: While the graph includes subject-bridge edges, deeper integrative questions (combining multiple subjects) are underrepresented.

16.2 Recommended Enhancements

(1) **A/B Testing Framework**: Implement experiment tracking so that recommendation engine variants can be systematically compared on engagement, learning gain, and score prediction accuracy.

(2) **Spaced Repetition Scheduler**: Integrate a Leitner-box style system for proactive weak-topic revisiting at optimal intervals.

(3) **Prerequisite Awareness**: Explicitly model prerequisites and enforce sequencing (don't serve eigenvalue questions until linear-solve fundamentals are mastered).

(4) **In-App Longitudinal Dashboards**: Display student ELO trends, weak-topic evolution, and predicted score trajectory directly in student and teacher dashboards.

(5) **Adaptive Difficulty Calibration**: Use item-response theory (IRT) to calibrate question difficulties and adapt them based on population performance.

(6) **Peer Benchmarking**: Compare student performance against class cohort (anonymous) to provide contextualized feedback.

(7) **Intelligent Content Recommendation**: Beyond questions, recommend:

- Reading materials or videos for weak topics.
- Worked examples and step-by-step solutions.
- Similar previously-solved questions for reinforcement.

(8) **Real-Time Collaboration Features**: Enable peer study groups, shared notes, and instructor office hours within the platform.

17 Conclusion

GATEWay is a comprehensive, multi-dimensional educational intelligence platform that transcends simple mock-test portals. Its core value lies not in any single component, but in the seamless integration of:

- **Graph-based adaptive sequencing** that personalizes question difficulty and topic selection.
- **Multiple specialized recommendation engines** addressing next-best questions, remediation, rapid-guess detection, and personalized coaching.
- **Classroom orchestration tools** enabling teachers to author assignments, monitor progress, and receive intervention recommendations.
- **Comprehensive telemetry** capturing both persistent state (ELO, scores) and event streams (actions, decisions) for deep observability.
- **Predictive scoring** powered by a rigorous, four-component formula that estimates likely GATE scores with confidence bands.
- **Rule-based coaching system** that delivers personalized study recommendations and guidance based on student performance analysis.
- **Explainable, replayable history** enabling students to revisit and learn from past attempts.
- **Research and analytics infrastructure** supporting longitudinal studies and exploratory insights.

This combination positions GATEWay as a strong candidate for:

- **Advanced undergraduate or graduate project submissions**: Demonstrating systems integration, machine learning, and educational technology depth.

Unpublished working draft. Not for distribution.

- **Technical job portfolios**: Showcasing full-stack development, architectural thinking, and domain expertise.
- **Academic research**: Publishing work on adaptive learning, recommendation systems, or educational analytics.
- **StartUp/Product pitches**: Presenting a differentiated solution addressing real learner and educator needs.
- **Viva and technical interviews**: Discussing complex systems, tradeoff decisions, and future research directions.

The platform demonstrates that educational software, when thoughtfully designed, can elevate not just test scores but learner autonomy, teacher effectiveness, and the quality of educational evidence.

Unpublished working draft.
Not for distribution.